

UCLA CS 251a: Advanced Computer Architecture

Lecture 2 : Technology & Performance

Prof. Tony Nowatzki

Slide History/Attribution Diagram:



How do we build a high-**performance** computer?

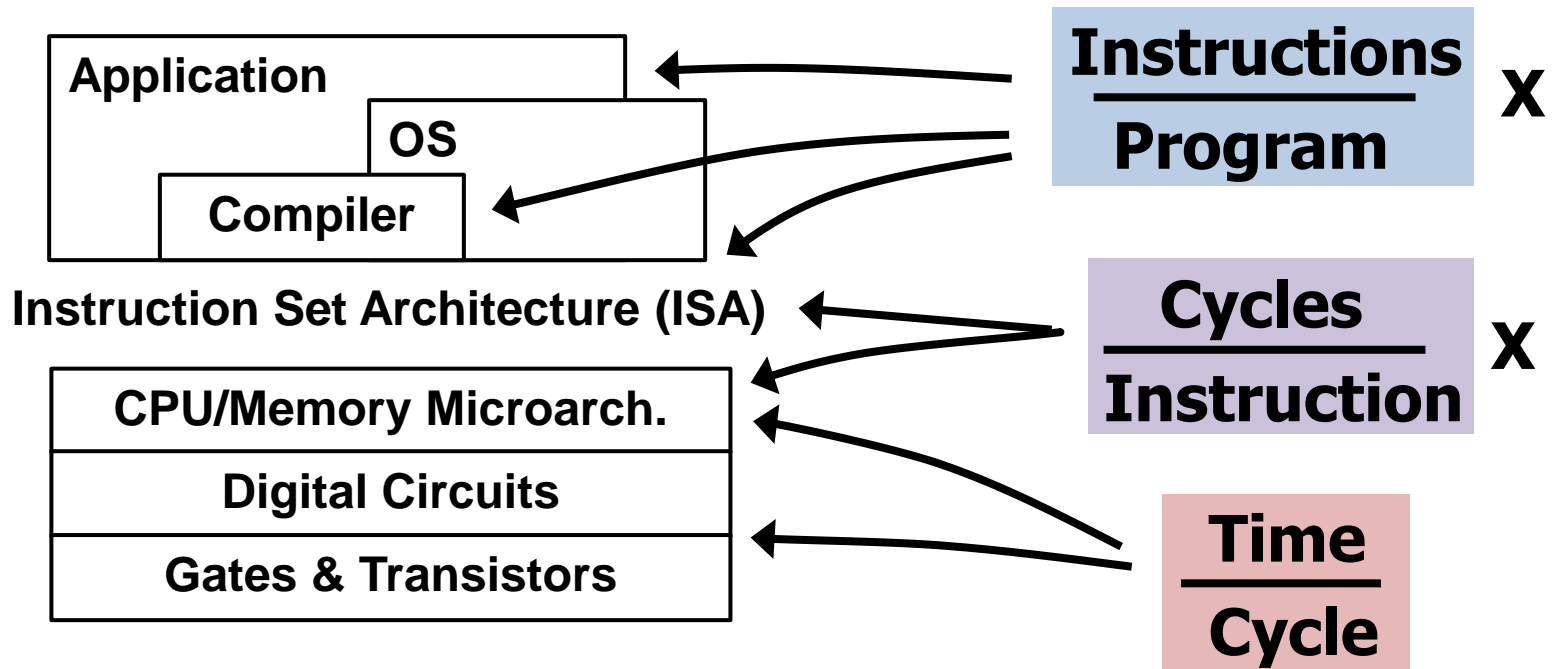
Subject to: Cost,
Power,
Energy,
Reliability,
Security,
Generality

Iron Law of Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Time}}{\text{Cycle}}$$

- Programs consist of simple operations (instructions)
 - Add two numbers, fetch data value from memory, etc.
- **Instructions per program**: “dynamic instruction count”
 - Runtime count of instructions executed by the program
 - Determined by program, compiler, instruction set architecture (ISA)
- **Cycles per instruction**: “CPI” (typical range: 4 to 1/4)
 - On average, how many *cycles* does an instruction take to execute?
 - Determined by program, compiler, ISA, micro-architecture
- **Seconds per cycle**: clock period, length of each cycle
 - Inverse metric: cycles per second (Hertz) or cycles per ns (Ghz)
 - Determined by micro-architecture, **technology parameters**

Computer System Layers





Main Focus for Today

Computer Building Blocks

- Transistors/Gates (computing)
 - How can they be connected to do something useful?
 - How do we evaluate how fast a logic block is?
- Wires (transporting)
 - What and where are they?
 - How can they be modeled?
- Memories (storing)
 - SRAM (speed) vs. DRAM (density)
- Compute and Store:
 - CAMs: Memories enabling associative lookups (lookup by value rather than by index)
 - Processing-in-memory
 - Digital: vectorized logical ops over some (2-3) RAM bitlines
 - Analog: arithmetic ops over many RAM bitlines
 - e.g. Memristors: Memories with programmable resistance


What are microprocessors made of?




Xeon Circa 2005

Semiconductor Technology Background

Transistor Circa 1947





Transistor Circa Now



Human hair $\sim 100\mu\text{m}$,
(10000 bigger than 10nm transistor)

Basic technology element: MOSFET



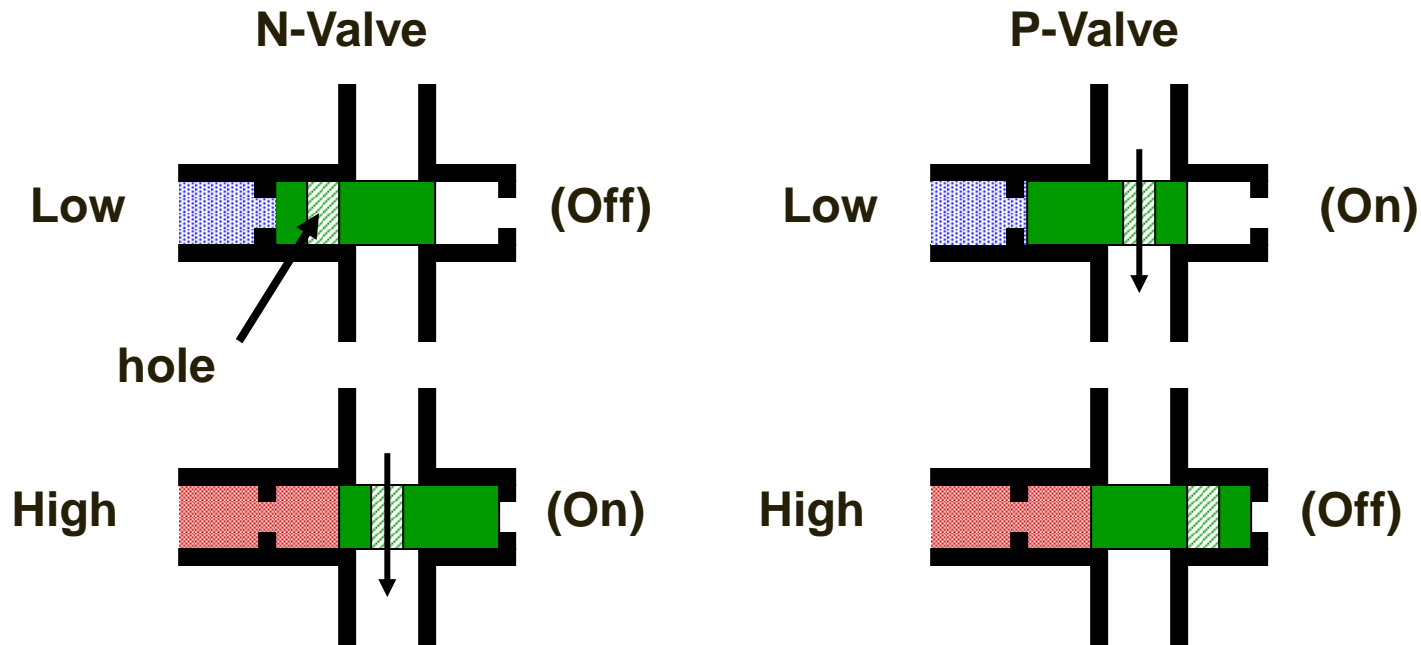
- Solid-state component acts like electrical switch
- **MOS**: three materials needed to make a transistor
 - **Metal**: Aluminum, Tungsten, Copper: conductor
 - **Oxide**: Silicon Dioxide (SiO_2): insulator
 - **Semiconductor**: doped Si: conducts under certain conditions
- **FET**: field-effect (the mechanism) transistor
 - Voltage on gate: current flows source to drain (transistor on)
 - No voltage on gate: no current (transistor off)
- **Channel length**: characteristic parameter (short \rightarrow fast)
 - Aka "feature size" or "technology"
 - Currently: 5-7 nm (but doesn't really measure anything anymore)
 - Continued miniaturization (scaling) known as "**Moore's Law**"

How do Transistors work?


An analogy for CS simpletons (i.e. me)
:)

A Transistor Analogy: Computing with Air


- Use air pressure to encode values
 - High pressure represents a "1" (blow)
 - Low pressure represents a "0" (suck)
- Valve can allow or disallow the flow of air
 - Two types of valves




Pressure Inverter




Pressure Inverter (Low to High)




Pressure Inverter



Pressure Inverter (High to Low)



What's that gate?




(nand)

Analogy Explained

- Pressure differential → electrical potential (voltage)
 - Air molecules → electrons
 - Pressure (molecules per volume) → voltage
 - High pressure → high voltage
 - Low pressure → low voltage
- Air flow → electrical current (eh, good enough)
 - Pipes → wires
 - Air only flows from high to low pressure
 - Electrons only flow from high to low voltage
 - Flow only occurs when changing from 1 to 0 or 0 to 1
- Valve → transistor
 - The transistor: one of the century's most important inventions

Transistors as Switches


- Two types
 - N-type: Conduct when gate voltage is 1
 - P-type: Conduct when gate voltage is 0
- Properties
 - Solid state (no moving parts)
 - Reliable (low failure rate)
 - Small (10nm channel length)
 - Fast (<0.1ns switch latency)
- **CMOS**: complementary n-/p-networks form boolean logic



CMOS Examples


- Example I: inverter

- Case I: input = 0
 - P-transistor closed, n-transistor open
 - Power charges output (1)
- Case II: input = 1
 - P-transistor open, n-transistor closed
 - Output discharges to ground (0)



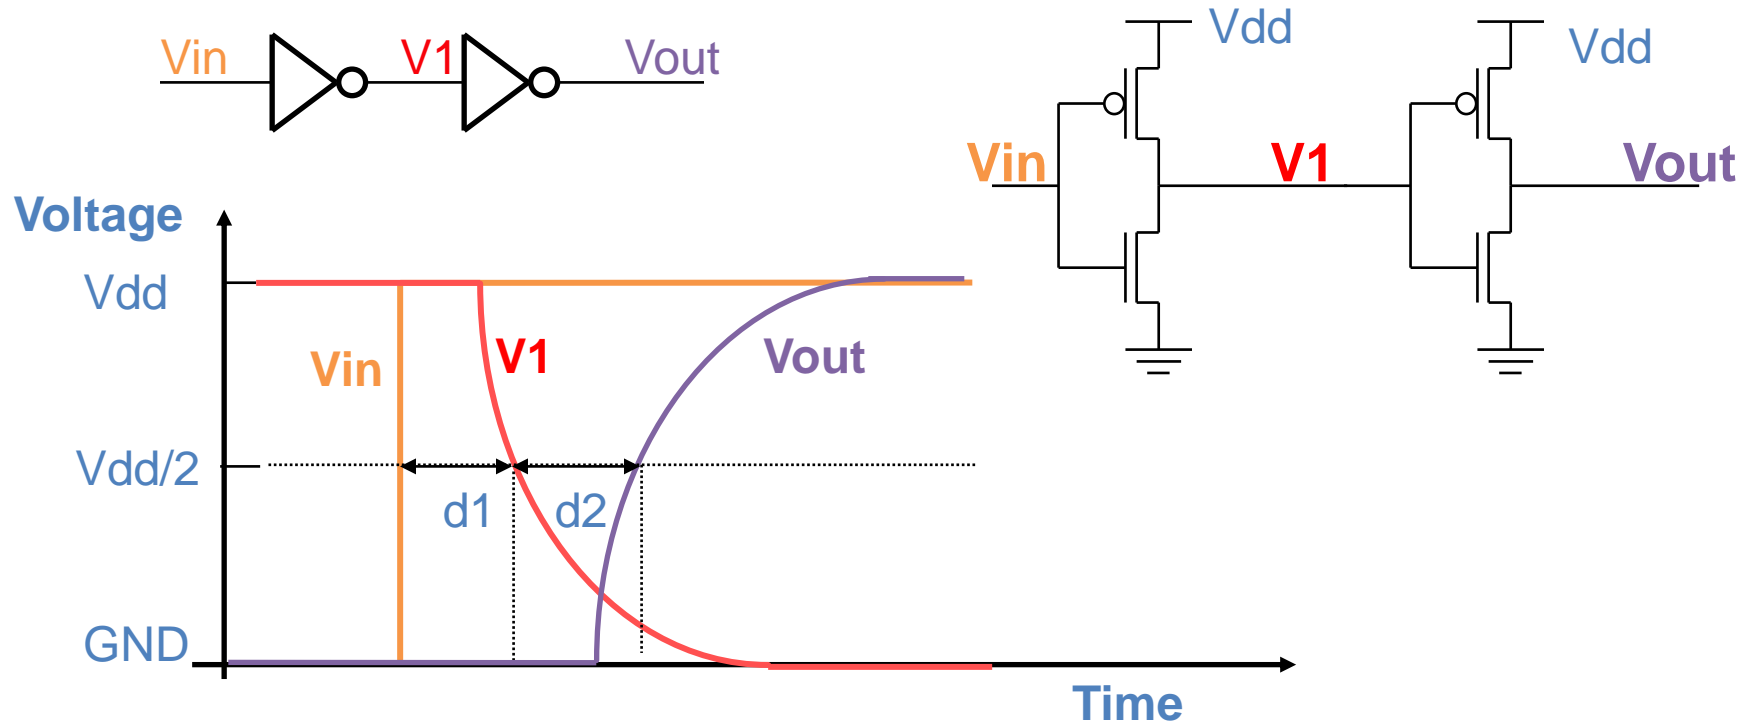
- Example II: look at **truth table**

- $0, 0 \rightarrow 1$ $0, 1 \rightarrow 1$
- $1, 0 \rightarrow 1$ $1, 1 \rightarrow 0$
- Result: this is a **NAND** (NOT AND)
- NAND is universal (can build any logic function)



Circuit Timing & Critical Path

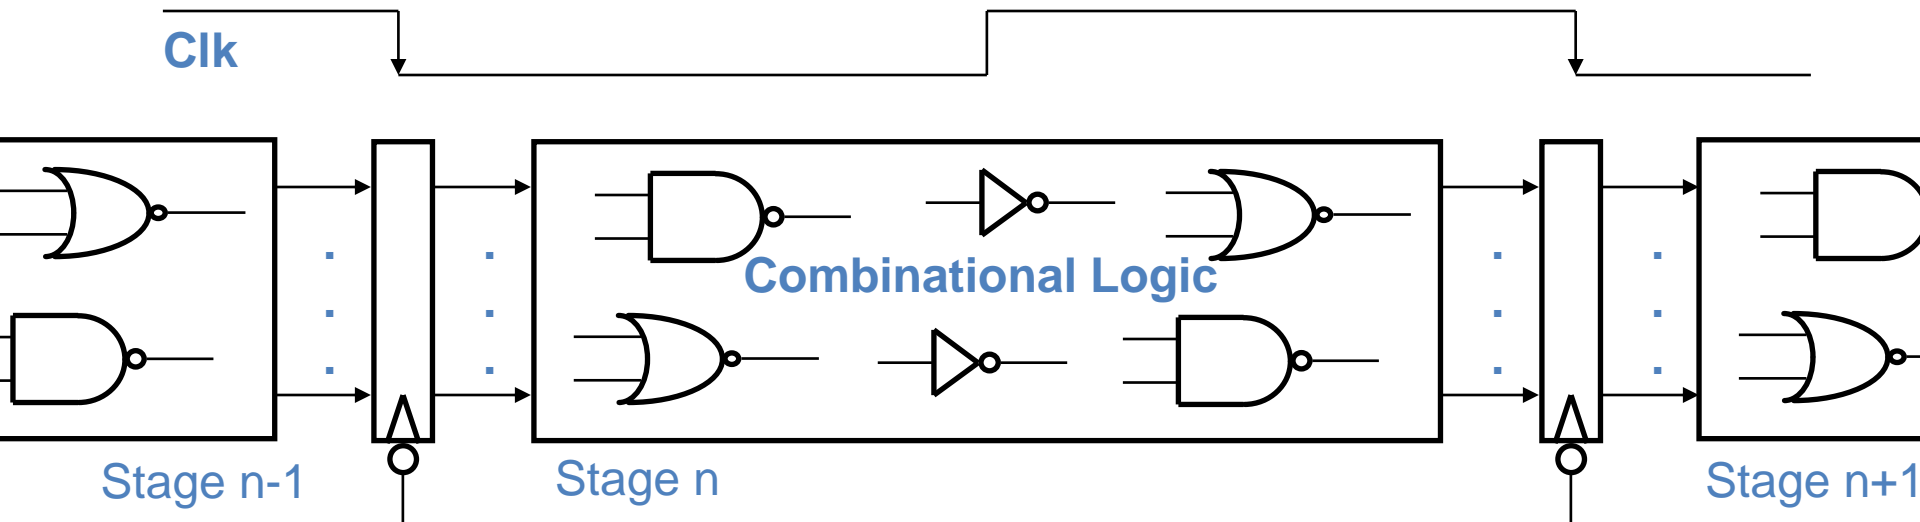
What about Transistor Delay



- Total Propagation Delay = Sum of individual delays = $d1 + d2$


Synchronous Design

- Hard to coordinate design components with logic only
 - (can be done, called asynchronous design)
- Synchronous: State elements synchronized by clock





- All storage elements are clocked by the same clock edge
- The combination logic block's:
 - Inputs are updated at each clock tick
 - All outputs **MUST** be stable before the next clock tick

Critical Path & Cycle Time





- Critical path: the slowest path between any two storage devices
- Cycle time $>$ time critical path
- **This is where frequency comes from :)**

Technology Basis of Transistor Speed



- Physics 101: delay through an electrical component $\sim RC$
 - **Resistance (R)**  $\sim \text{length} / \text{cross-section area}$
 - Slows rate of charge flow
 - **Capacitance (C)**  $\sim \text{length} * \text{area} / \text{distance-to-other-plate}$
 - Stores charge
 - **Voltage (V)**
 - Electrical pressure
 - **Threshold Voltage (V_t)**
 - Voltage at which a transistor turns "on"
 - Property of transistor based on fabrication technology
 - **Switching time $\sim t_o (R * C) / (V - V_t)$**
- Two kinds of electrical components
 - CMOS transistors (gates)
 - Wires

Low V_t (faster, more power)



High V_t (slower, less power)

Wire Geometry



IBM CMOS7, 6 layers of copper wiring

- Wires 4-dimensional: **length, width, height, "pitch"**
 - Longer wires have more resistance
 - "Thinner" wires have more resistance
 - Closer wire spacing ("pitch") increases capacitance

Increasing Problem: Wire Delay

- RC Delay of wires
 - **Resistance** proportional to: $\text{resistivity} * \text{length} / (\text{cross section})$
 - Wires with smaller cross section have higher resistance
 - Resistivity (type of metal, copper vs aluminum)
 - **Capacitance** proportional to length
 - And wire spacing (closer wires have large capacitance)
 - Permittivity or “dielectric constant” (of material between wires)
- Result: delay of a wire is **quadratic** in length
 - Insert “inverter” repeaters for long wires
 - Why? To bring it back to linear delay... but repeaters still add delay
- Trend: wires are getting relatively slow to transistors
 - And relatively longer time to cross relatively larger chips
- **Arch. Implication:**
 - **Avoid limit the amount of communication**
 - **Keep high-bandwidth communication as local as possible**

Power & Energy

Power/Energy Are Increasingly Important


- **Battery life** for mobile devices
 - Laptops, phones, cameras
- **Tolerable temperature** for devices without active cooling
 - Power means temperature, active cooling means **cost**
 - No room for a fan in a cell phone, no market for a hot cell phone
- **Electric bill** for compute/data centers
 - Pay for power twice: once in, once out (to cool)
- **Environmental concerns**
 - “Computers” account for growing fraction of energy consumption
 - As a percentage of global carbon emissions
 - Bitcoin: ~0.1%
 - Smartphones: ~0.3%
 - Datacenters: ~3% (caveat: a lot of it is manufacturing)
 - Commercial Flights: ~2.4%

Energy & Power

- **Energy**: measured in Joules or Watt-seconds
 - Total amount of energy stored/used
 - Battery life, electric bill, environmental impact
 - Joules per Instruction (car analogy: gallons per mile)
- **Power**: energy per unit time (measured in Watts)
 - Joules per second (car analogy: gallons per hour)
 - Power impacts power supply and cooling requirements (cost)
 - Power-density (Watt/mm^2): important related metric
 - Peak power vs average power
 - E.g., camera, power “spikes” when you actually take a picture
- Two sources:
 - **Dynamic power**: active switching of transistors
 - **Static power**: leakage of transistors even while inactive

Dynamic Power

- **Dynamic power (P_{dynamic}):** aka switching or active power
 - Energy to switch a gate (0 to 1, 1 to 0)
 - Each gate has capacitance (C)
 - Energy to charge/discharge a capacitor is \sim to $C * V^2$
- **$P_{\text{dynamic}} \sim N * C * V^2 * f * A$**
 - N: number of transistors
 - C: capacitance per transistor (size of transistors)
 - V: voltage (supply voltage for gate)
 - f: frequency (transistor switching freq. is \sim to clock freq.)
 - A: activity factor (not all transistors may switch this cycle)




Reducing Dynamic Power

- Target each component: $P_{\text{dynamic}} \sim N * C * V^2 * f * A$
- **Reduce number of transistors** (N)
 - Use fewer transistors/gates (better design; specialized hardware)
- **Reduce capacitance** (C)
 - Smaller transistors (Moore's law)
- **Reduce voltage** (V)
 - Quadratic reduction in energy consumption!
 - But also slows transistors (recall transistor speed is \sim to V)
- **Reduce frequency** (f)
 - Slower clock frequency (reduces power but not energy) Why?
- **Reduce activity** (A)
 - "Clock gating" disable clocks to unused parts of chip
 - Don't switch gates unnecessarily

Static Power

- **Static power (P_{static})**: aka idle or leakage power
 - Transistors don't turn off all the way
 - Transistors "leak"
 - Analogy: leaky valve
 - $P_{\text{static}} \sim N * V * e^{-V_t}$
 - N: number of transistors
 - V: voltage
 - **V_t (threshold voltage)**: voltage at which transistor conducts (begins to switch)
- Switching speed vs leakage trade-off
- The lower the **V_t** :
 - Faster transistors (linear)
 - Transistor speed $\sim V - V_t$
 - Leakier transistors (exponential)



Reducing Static Power

- Target each component: $P_{\text{static}} \sim N * V * e^{-V_t}$
- **Reduce number of transistors** (N)
 - Use fewer transistors/gates
- **Disable transistors** (also targets N)
 - “Power gating” disable power to unused parts (long latency to power up)
 - Power down units (or entire cores) not being used
- **Reduce voltage** (V)
 - Linear reduction in static energy consumption
 - But also slows transistors (transistor speed is \sim to V)
- **Dual V_t** – use a mixture of high and low V_t transistors
 - Use slow, low-leak transistors in SRAM arrays
 - Requires extra fabrication steps (cost)
- Note: reducing frequency can actually hurt static energy. Why?

Scale Voltage & Frequency Together

- Recall: $P_{\text{dynamic}} \sim N * C * V^2 * f * A$
 - Because max frequency \sim to V ...
 - $P_{\text{dynamic}} \sim$ to V^3
- Reduce both voltage and frequency linearly
 - **Cubic decrease in dynamic power**
 - Linear decrease in performance (sub-linear if memory bound)
 - Thus, only about quadratic in energy
 - Linear decrease in static power
 - Thus, only modest static energy improvement
- Example: Intel Xscale
 - 1 GHz \rightarrow 200 MHz reduces energy used by 30x
 - But around 5x slower
 - 5 x 200 MHz in parallel, use **1/6th the energy**
- **Power-limited designs favor multi-cores!**

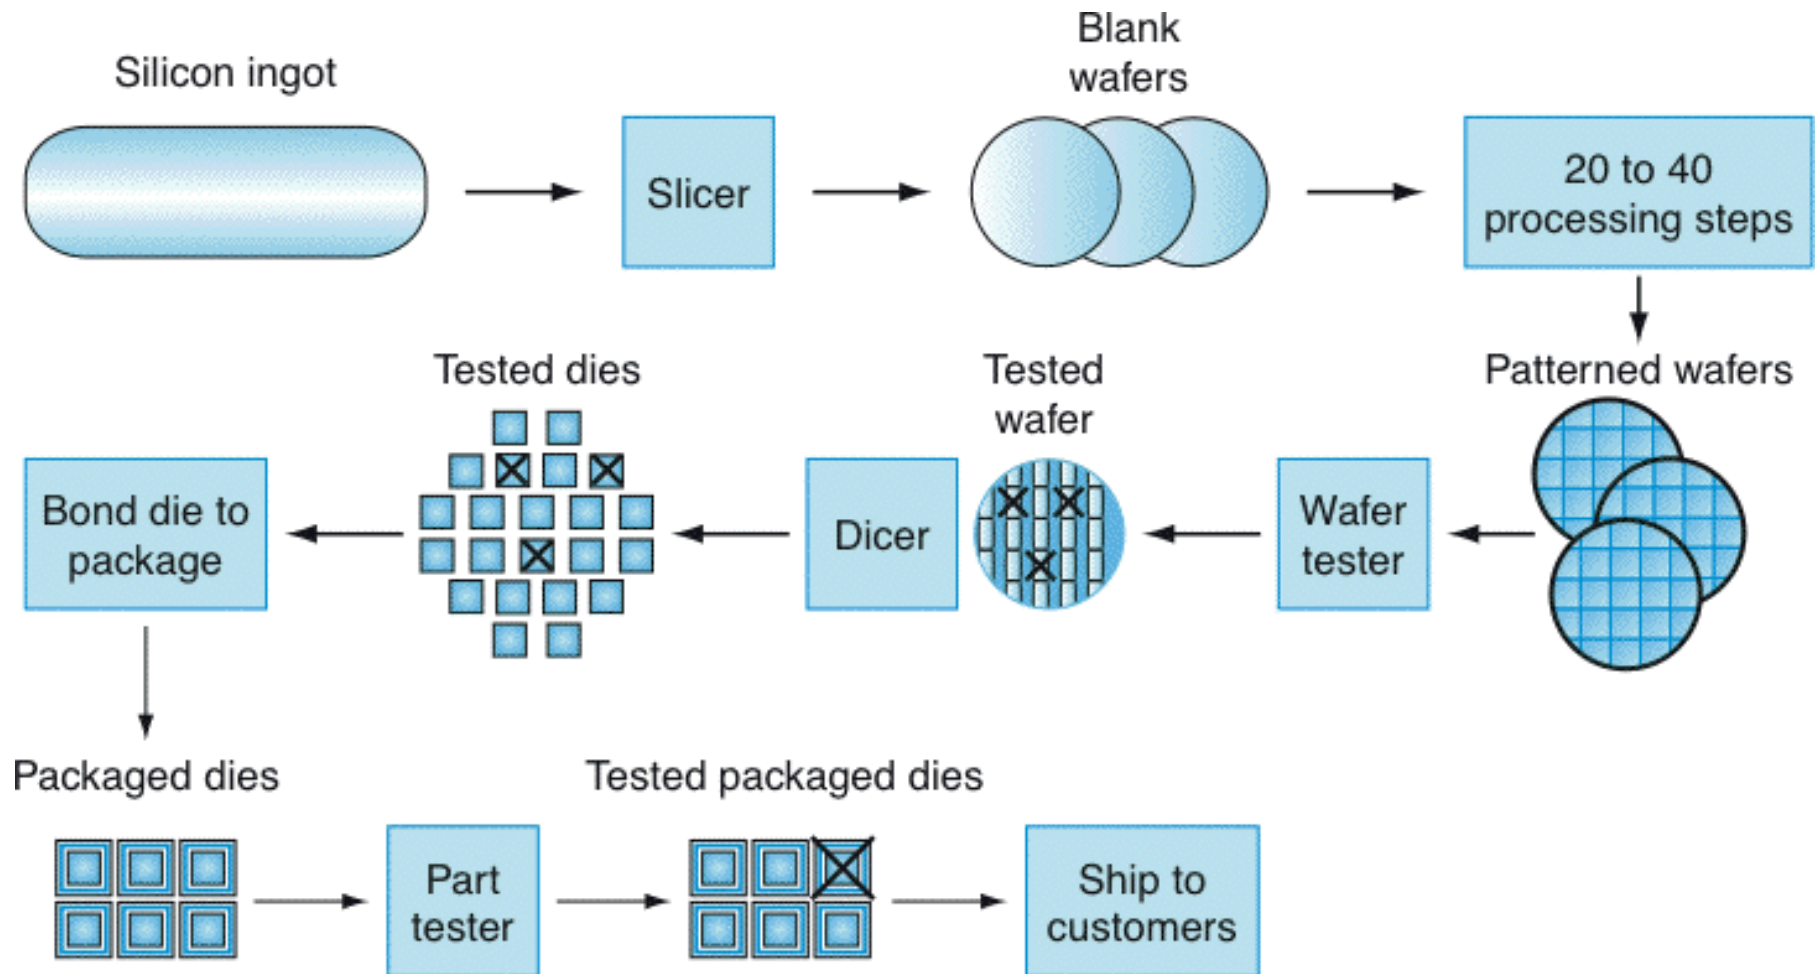
Dynamic Voltage/Frequency Scaling

- **Dynamically trade-off performance for lower power**
 - Change the voltage and frequency at runtime
 - Under control of operating system
 - Best of both worlds?
- Modern chips can adjust frequency on a per-core basis

	Mobile PentiumIII "SpeedStep"	Transmeta 5400 "LongRun"	Intel X-Scale (StrongARM2)
f (MHz)	300–1000 (step=50)	200–700 (step=33)	50–800 (step=50)
V (V)	0.9–1.7 (step=0.1)	1.1–1.6V (cont)	0.7–1.65 (cont)
High-speed	3400MIPS @ 34W	1600MIPS @ 2W	800MIPS @ 0.9W
Low-power	1100MIPS @ 4.5W	300MIPS @ 0.25W	62MIPS @ 0.01W


Manufacturing:
Get those transistors on a chip!

Manufacturing Steps

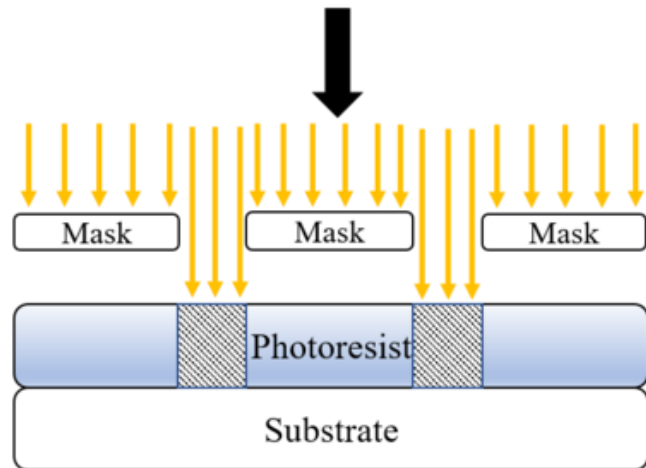


Lithography Step


(1) Apply photoresist




(2) Expose to light



Photomask




Negative Photoresist




Manufacturing Steps

- Multi-step photo-/electro-chemical process
 - More steps, higher unit cost
- + Fixed cost mass production (\$1 million+ for “mask set”)




Manufacturing Defects


Correct:




Defective:



Defective:




Slow:




- Defects can arise
 - Under-/over-doping
 - Over-/under-dissolved insulator
 - Mask mis-alignment
 - Particle contaminants
- Try to minimize defects
 - Process margins
 - Design rules
 - Minimal transistor size, separation
- Or, tolerate defects
 - Redundant or “spare” memory cells
 - Can substantially improve yield

Cost Implications of Defects

- Cost of manufacturing process depends on:
 - Size of the wafer
 - Number of steps
- Cost/mm² of a **chip** depends on *chip area*
 - Why? random defects
 - Larger chip: more chance of defect



- **Wafer yield**: % wafer that is chips
- **Die yield**: % chips that work
- Yield is increasingly non-binary – fast vs slow chips (instead of working vs broken)



All Roads Lead To Multi-Core

+ Multi-cores reduce unit costs

- Higher yield than same-area single-cores
- Why? Defect on one of the cores? Sell remaining cores for less
- This is part of why you see so many versions of CPUs or GPUs in a single generation with slightly different core counts

+ Multi-cores can reduce design costs too


- Replicate existing designs rather than re-design larger single-cores

Fixed Costs

- For new chip design
 - Design & verification: ~\$100M (500 person-years @ \$200K per)
 - Amortized over “proliferations”, e.g., Core i3, i5, i7 variants
- For new (smaller) technology generation
 - ~\$3B for a new fab (future fabs probably a lot more)
 - Amortized over multiple designs (Intel’s tick/tock)
 - Amortized by “rent” from companies that don’t fab themselves
 - Xilinx invented this approach
 - Eg. Qualcomm, Apple, NVIDIA, Broadcom, MediaTek, AMD
- Moore’s Law generally increases startup cost
 - More expensive fabrication equipment
 - More complex chips take longer to design and verify

Technology Scaling

Moore's Law: Technology Scaling



- **Moore's Law:** aka "technology scaling"
 - Continued miniaturization (esp. reduction in channel length)
 - + Improves switching speed, area(cost)/transistor
 - Reduces transistor reliability
 - eg: DRAM density (transistors/area) doubles every 18 months
- Dennard Scaling: power/transistor
 - Corollary to Moore's law that has more-or-less ended

Moore's Effect #1: Transistor Count

- Linear shrink in each dimension
 - 180nm, 130nm, 90nm, 65nm, 45nm, 32nm, ...
 - Each generation is a 1.414 linear shrink
 - Shrink each dimension (2D)
 - Results in 2x more transistors (1.414×1.414)
- Reduces cost per transistor
- More transistors can increase performance
 - Job of a computer architect: use the ever-increasing number of transistors
 - Examples: caches, exploiting parallelism at all levels

Moore's Effect #2: RC Delay

- **First-order: speed scales proportional to gate length**
 - Has provided much of the performance gains in the past
- Scaling helps wire and gate delays in some ways...
 - + Transistors become shorter (Resistance↓), narrower (Capacitance↓)
 - + Wires become shorter (Length↓ → Resistance↓)
 - + Wire "surface areas" become smaller (Capacitance↓)
- Hurts in others...
 - Transistors become narrower (Resistance↑)
 - Gate insulator thickness becomes smaller (Capacitance↑)
 - Wires becomes thinner (Resistance↑)
- What to do?
 - Take the good, use wire/transistor sizing & repeaters to counter bad
 - Keep Communication local!
 - Exploit new materials: Aluminum → Copper, metal gate, high-K
- **Used to get much faster, not as much any more...**

Moore's Effect #3: Cost

- Mixed impact on unit integrated circuit cost
 - + Either lower cost for same functionality...
 - + Or same cost for more functionality
 - Difficult to achieve high yields (defects)
 - Process variation
- Increases startup cost
 - More expensive fabrication equipment
 - Takes longer to design, verify, and test chips
 - Being on the cutting edge is expensive!

Moore's Effect #4: Power

- + Technology scaling reduces power/transistor...
 - Reduced sizes and surface areas reduce capacitance (C)
 - But this effect is small now (Dennard scaling is over)
- ...but increases power density and total power
 - By increasing transistors/area and total transistors
 - Faster transistors → higher frequency → more power
 - Hotter transistors leak more (thermal runaway)
- The end of voltage scaling & “dark silicon”
 - Dark silicon: can't use all of your transistors at the same time with maximum frequency
 - Two ways to deal with this:
 1. Lower frequency, do more in parallel
 2. More specialized units that you only operate some of the time



Trends in Power

	386	Pentium	Pentium II	Pentium 4	Core2	Core i7	Cascade Lake
Year	1985	1993	1998	2001	2006	2009	2019
Technode (nm)	1500	350	180	130	65	45	14
Transistors (M)	0.3	3.1	5.5	42	291	731	8000
Voltage (V)	5	3.3	2.9	1.7	1.3	1.2	1.2
Clock (MHz)	16	66	200	1500	3000	3300	2.6-3.8K
Power (W)	1	16	35	80	75	130	100-400
Peak MIPS	6	132	600	4500	24000	52,800	582,400
MIPS/W	6	8	17	56	320	406	1,456

- Supply voltage decreasing over time
 - But “voltage scaling” is perhaps reaching its limits
- Emphasis on power starting around 2000
 - Resulting in slower frequency increases
 - Also note number of cores increasing (4 in Core i7, 28 for Casc. Lake)

Processor Power Breakdown

- Power breakdown for IBM POWER4
 - Two 4-way superscalar, 2-way multi-threaded cores, 1.5MB L2
 - Big power components are L2, data cache, scheduler, clock, I/O
 - Implications on “complicated” versus “simple” cores



Moore's Effect #5: Psychological / Economic


- **Moore's Curve:** common interpretation of Moore's Law
 - "CPU performance doubles every 18 months"
 - Self fulfilling prophecy: 2X every 18 months is $\sim 1\%$ per week
 - Q: Would you add a feature that improved performance 20% if it would delay the chip 8 months?
 - Processors under Moore's Curve (arrive too late) fail spectacularly
 - E.g., Intel's Itanium, Sun's Millennium

Moore's Law in the Future

- Won't last forever, approaching physical limits
 - "If something must eventually stop, it can't go on forever"
 - But betting against it has proved foolish in the past
 - Perhaps will "slow" rather than stop abruptly
- Transistor count will likely continue to scale in short term
 - "Die stacking" is becoming main stream
 - Uses the third dimension to increase transistor count (mostly for integrating logic/memory, or stacking memory -- HBM)
 - Wafer-scale computing – don't use packages, just use wafer
 - Chiplets/Dielets – tightly integrate many small dies
- But transistor performance scaling?
 - Running into physical limits
 - Example: gate oxide is less than 10 silicon atoms thick!
 - Can't decrease it much further
 - Power is becoming a limiting factor

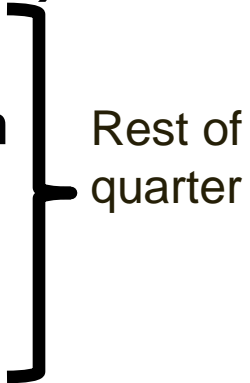
Faith no Moore

Selected predictions for the end of Moore's law



Sources: Intel; press reports; *The Economist*

Technology Summary

- Has a first-order impact on computer architecture
 - Cost (die area)
 - Performance (transistor delay, wire delay)
 - **Changing rapidly**
 - Most significant trends for architects (and thus 251a)
 - More and more transistors
 - What to do with them? → integration → **parallelism**
 - Logic is improving faster than memory & cross-chip wires
 - “Memory wall” → caches, more integration
 - Techniques to localize communication
- 
- Rest of quarter

Performance

Iron Law of Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Time}}{\text{Cycle}}$$

CPI
(microarchitecture value added)

Performance

- Two definitions
 - **Latency (execution time)**: time to finish a fixed task
 - **Throughput (bandwidth)**: number of tasks in fixed time
 - Very different: throughput can exploit parallelism, latency cannot
 - Ex: baking bread
 - Often contradictory
 - Choose definition of performance that matches your goals
- Example: move people from A \rightarrow B, 10 miles
 - Car: capacity = 5, speed = 60 miles/hour
 - Bus: capacity = 60, speed = 20 miles/hour
 - Latency: **car = 10 min**, bus = 30 min
 - Throughput: car = 15 PPH (count return trip), **bus = 60 PPH**

Performance Improvement

- Given a program of interest P...
- Speedup Ratio of Processor A over B is either:
 - $\text{Latency}(P, B) / \text{Latency}(P, A)$
 - $\text{Throughput}(P, A) / \text{Throughput}(P, B)$
- Speedup Percent of Processor A over B is either:
 - $(\text{Latency}(P, B) / \text{Latency}(P, A) - 1) * 100$
 - $(\text{Throughput}(P, A) / \text{Throughput}(P, B) - 1) * 100$
- Car/bus example
 - Latency? Car is 3 times faster than bus
 - So is it 200% or 300% faster?
- What is 'P' in $\text{Latency}(P, A)$?
 - $\text{Latency}(A)$ makes no sense, processor executes **some** program

Adding/Average Performance Numbers

- You can add latencies, but not throughput
 - $\text{Latency}(P1+P2, A) = \text{Latency}(P1, A) + \text{Latency}(P2, A)$
 - $\text{Throughput}(P1+P2, A) \neq \text{Throughput}(P1, A) + \text{Throughput}(P2, A)$
 - $\text{Throughput}(P1+P2, A) = 2 / [(1/\text{Throughput}(P1, A)) + (1/\text{Throughput}(P2, A))]$
 - E.g. 1 mile @ 30 miles/hour + 1 mile @ 90 miles per hour
 - Average is **not** 60 miles per hour
 - 0.03333 hours at 30 miles/hour + 0.01111 hours at 90 miles/hour
 - Average is only 45 miles/hour! (2 miles/ (0.03333 + 0.01111))
- In general for means (average):
 - **Arithmetic**: $(1/N) * \sum_{P=1..N} \text{Latency}(P)$
 - For units that are proportional to time (e.g., latency)
 - **Harmonic**: $N / \sum_{P=1..N} 1/\text{Throughput}(P)$
 - For units that are inversely proportional to time (e.g., throughput)
 - **Geometric**: $\sqrt[N]{\prod_{P=1..N} \text{Speedup}(P)}$
 - For unitless quantities (e.g., speedup ratios)
 - Only option for ratios, as geomean is the only one where the relative merit does not depend on the reference/baseline speedup

Pitfalls of Partial Performance Metrics

Danger: Partial Performance Metrics

- General public often equates performance with particular architecture aspects, **neglecting CPI!**
 - ISA Bitwidth (16,32,64,128 bit, yay!)
 - **Frequency**
 - # Cores
- Which processor would you buy?
 - Processor A: CPI = 2, clock = 5 GHz
 - Processor B: CPI = 1, clock = 3 GHz
 - Probably A, but B is faster (assuming same ISA/compiler)
- Classic example
 - 800 MHz PentiumIII faster than 1 GHz Pentium4!
- **Meta-point: danger of partial performance metrics!**

MIPS Metric (ignores instr. count)

- (Micro) architects often ignore dynamic instruction count
 - Typically work in one ISA/one compiler → treat it as fixed

$$\frac{\text{Time}}{\text{Instruction}} = \cancel{\frac{\text{Instructions}}{\text{Program}}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Time}}{\text{Cycle}}$$

- **MIPS** (millions of instructions per second, *inverse of above*)
 - **Instructions/second * 10⁻⁶**
 - **Cycles / second**: clock frequency (in MHz)
 - Example: CPI = 2, clock = 500 MHz → 0.5 * 500 MHz = 250 MIPS
- Pitfall of MIPS: may vary inversely with actual performance
 - Compiler removes insns, program gets faster, MIPS can go down (eg. instructions that it removed were easy to run in parallel...)
- Other problems:
 - Some optimizations actually add instructions
 - Work per instruction varies (e.g., multiply vs. add, FP vs. integer)
 - ISAs are not equivalent

MFLOPS (MegaFLOPS)

- **MFLOPS**: like MIPS, but counts only FP ops, because...
 - + FP ops can't be optimized away (by compiler)
 - + FP ops have longest latencies (FP Units are larger)
 - + FP ops are same across machines
- Pitfalls:
 - Many CPU programs are “integer” – light on FP
 - Loads from memory can take much longer than even an FP divide
 - Even FP instruction sets are not equivalent
- Upshot: Neither MIPS nor MFLOPS are universal

Cycles per Instruction (CPI)

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Time}}{\text{Cycle}}$$

- This course is mostly about improving **CPI**
 - Cycle/instruction for average instruction
 - $\text{IPC} = 1/\text{CPI}$
 - Used more frequently than CPI, but harder to compute with
 - Different instructions have different cycle costs
 - E.g., integer add typically takes 1 cycle, FP divide takes > 10
 - Assumes you know something about instruction frequencies
- CPI example
 - A program executes equal integer, FP, and memory operations
 - Cycles per instruction type: integer = 1, memory = 2, FP = 3
 - What is the CPI? $(0.33 * 1) + (0.33 * 2) + (0.33 * 3) = 2$
 - **Caveat:** this sort of calculation ignores dependencies completely
 - Kind-of useful for back-of-the-envelope arguments?

Improving CPI

- This course focuses on improve CPI instead of frequency
 - Until early 2000s, clock accounts for 70%+ of performance improvement
 - Achieved via deeper pipelines
 - That has been forced to change
 - Deep pipelining is **not** power efficient
 - Physical speed limits are approaching
 - 1 GHz: 1999, 2 GHz: 2001, 3 GHz: 2002, 3.8 GHz: 2004, 5 GHz: 2008
 - Intel Core 2: 1.8 – 3.2 GHz in 2008
 - **Techniques we'll look at:**
 - **Caching, speculation, multiple issue, out-of-order issue**
 - **Vectors, multiprocessing, more ...**
- Moore helps because CPI reduction requires transistors
 - The definition of parallelism is “more transistors”
 - Simplest example is caches

Performance Rules of Thumb



- Make common case fast
 - **"Amdahl's Law"**
 - $\text{Time}_{\text{optimized}} = \frac{\text{Time}_{\text{orig}} * \text{fraction}_x}{\text{Speedup}_x} + \text{Time}_{\text{orig}} * (1 - \text{fraction}_x)$
 - $\text{Speedup} = \text{Time}_{\text{orig}} / \text{Time}_{\text{optimized}}$
 - $\text{Speedup} = 1 / ((1 - \text{fraction}_x) + \text{fraction}_x / \text{Speedup}_x)$
 - Corollary: don't optimize 5% to the detriment of the other 95%
 - $\text{Speedup}_{\text{overall}} = 1 / ((1 - 5\%) + 5\% / \infty) \approx 1.05$
- Build a balanced system
 - Don't over-engineer capabilities that cannot be utilized
 - Try to be "bound" by the most expensive resource (if not everywhere)
- Design for actual, not peak performance
 - For actual performance X, machine capability must be $> X$

Next Time...

- ISAs

Hidden Bonus Slides



Transistor Geometry: Width



Diagrams © Krste Asanovic, MIT

- **Transistor width**, set by designer for each transistor
- Wider transistors:
 - **Lower resistance** of channel (increases drive strength) – good!
 - But, **increases capacitance** of gate/source/drain – bad!
- Result: set width to balance these conflicting effects


Transistor Geometry: Length & Scaling



Diagrams © Krste Asanovic, MIT

- **Transistor length:** characteristic of “process generation”
 - 45nm refers to the transistor gate length, same for all transistors
- Shrink transistor length:
 - Lower resistance of channel (shorter) – good!
 - Lower gate/source/drain capacitance – good!
- Result: switching speed improves linearly as gate length shrinks

Amdahl's Law Graph



Source: Wikipedia

Little's Law

- Key Relationship between latency and bandwidth:
 - Average number in system = arrival rate * avg holding time
- Example:
 - How big a wine cellar should I build?
 - My family drinks (and buys) an average of 4 bottles per week
 - On average, I want to age my wine 5 years
 - # bottles in cellar = 4 bottles/week * 52 weeks/year * 5 years
= 1040 bottles (!)

More Little's Law

- How many outstanding cache misses?
 - Want to sustain 5 GB/s bandwidth
 - 64 byte blocks
 - 100ns miss latency
- Requests in system = arrival rate * time in system
$$= (5 \text{ GB/s} / 64 \text{ byte blocks}) * 100 \text{ ns}$$
$$= 8 \text{ misses}$$
- That's an **AVERAGE**.
 - Need to support many more if we hope to sustain this bandwidth
 - Rule of thumb: 2X

Reliability

Technology Basis for Reliability

- As transistors get smaller, they are less reliable
 - Wasn't a problem a few years ago, becoming a big problem
 - Small capacitance means fewer electrons represent 1 or 0
- **Transient faults**
 - A bit "flips" randomly, **temporarily**
 - Cosmic rays and such (more common at higher altitudes!)
 - Memory cells (especially memory) vulnerable today, logic soon
- **Permanent (hard) faults**
 - A gate or memory cell wears out, **breaks and stays broken**
 - Temperature & electromigration gradually deform components
- Solution for both: use **redundancy** to detect and tolerate

Aside: Memory Technology Families

- **SRAM:** “static” RAM
 - Used on processor chips (same transistors as used for “logic”)
 - Storage implemented as 6 transistors per bit
 - An inverter pair (2 transistors each) + two control transistors
 - Optimized for speed first, then secondarily density and power
- **DRAM (volatile memory):** “dynamic” RAM
 - Different manufacturing steps, not typically used on processor chips
 - Storage implemented as one capacitor + 1 transistor per bit
 - Optimized for density and cost
- **Flash (non-volatile memory):**
 - Used for solid state storage
 - Slower than DRAM, but non-volatile
- Disk is also a “technology”, but isn’t transistor-based

Memory Error Detection

- Idea: add extra state to memory to detect a bit flip
- **Parity**: simplest scheme
 - One extra bit, detects any single bit flip
 - Parity bit = $\text{XOR}(\text{data}_{N-1}, \dots, \text{data}_1, \text{data}_0)$
- Example:
 - 010101 $0 \wedge 1 \wedge 0 \wedge 1 \wedge 0 \wedge 1 = "1"$ so parity is "odd" (versus "even")
 - So, store "010101 **1**" in memory
 - When you read the data, and re-calculate the parity, say
 - 01**1**101 **1**, if the parity bit doesn't match, error detected
- Multiple bit errors? more redundancy can detect more

Memory Error Detection

- What to do on a parity error?
- **Crash**
 - **Dead programs tell no lies**
 - Fail-stop is better than silent data corruption
 - Avoiding writing that “\$1m check”
- For user-level data, OS can kill just the program
 - Not the whole system, unless it was OS data
- Alternative: correct the error

SEC Error Correction Code (ECC)

- **SEC**: single-error correct (a hamming code)
- Example: Four data bits, three "code" bits
 - $d_1 d_2 d_3 d_4 \mathbf{c_1 c_2 c_3} \rightarrow \mathbf{c_1 c_2} d_1 \mathbf{c_3} d_2 d_3 d_4$
 - $c_1 = d_1 \wedge d_2 \wedge d_4$, $c_2 = d_1 \wedge d_3 \wedge d_4$, $c_3 = d_2 \wedge d_3 \wedge d_4$
 - Syndrome: $c_i \wedge c'_i = 0$? no error
 - Otherwise, then $c'_3 c'_2 c'_1$ points to flipped-bit
- Working example
 - Original data = 0110 $\rightarrow c_1 = 1, c_2 = 1, c_3 = 0$
 - Flip $d_2 = 0010 \rightarrow c'_1 = 0, c'_2 = 1, c'_3 = 1$
 - Syndrome = 101 (binary 5) \rightarrow 5th bit? D_2
 - Flip $c_2 \rightarrow c'_1 = 1, c'_2 = 0, c'_3 = 0$
 - Syndrome = 010 (binary 2) \rightarrow 2nd bit? c_2

SECDED Error Correction Code (ECC)

- **SECDED**: single error correct, double error detect
- Example: $D = 4 \rightarrow C = 4$
 - $d_1 d_2 d_3 d_4 \mathbf{c_1 c_2 c_3} \rightarrow \mathbf{c_1 c_2} d_1 \mathbf{c_3} d_2 d_3 d_4 \mathbf{c_4}$
 - $c_4 = c_1 \wedge c_2 \wedge d_1 \wedge c_3 \wedge d_2 \wedge d_3 \wedge d_4$
 - Syndrome == 0 and $c'_4 == c_4 \rightarrow$ no error
 - Syndrome != 0 and $c'_4 != c_4 \rightarrow$ 1-bit error
 - Syndrome != 0 and $c'_4 == c_4 \rightarrow$ 2-bit error
 - Syndrome == 0 and $c'_4 != c_4 \rightarrow c_4$ error
 - **In general: $C = \log_2 D + 2$**
- Many machines today use 64-bit SECDED code
 - $C = 8$ (64bits + 8bits = 72bits, 12% overhead)
 - ChipKill - correct any aligned 4-bit error
 - If an entire memory chips dies, the system still works!

Another Issue: Process Variability

- As transistors get smaller...
 - Small geometric variations have relatively larger impact
- Example: Gate oxide thickness
 - In Intel's 65nm process: only 1.2 nm, just a few molecules thick!
 - Small variation in gate oxide thickness impacts speed and energy
 - Too thick: slow transistor
 - Too thin: exponential increase in leakage (static power)
- Some parts of the chip slow than others (impacts yield)
- Complicates high-speed memory designs
- Limits circuit techniques ("dynamic" versus "static" circuits)
 - Intel's Nehalem (Core i7) moved to all static circuits

Reliability

- **Mean Time Between Failures (MTBF)**
 - How long before you have to reboot or buy a new one
- CPU reliability is small in the grand scheme
 - Software most unreliable component in a system
 - Much more difficult to specify & test
 - Much more of it
 - Most unreliable hardware component: disk
 - Subject to mechanical wear

Moore's Bad Effect on Reliability

- Wasn't a problem until 5-10 years ago...
 - Except for transient-errors on chips in orbit (satellites)
 - ...a problem already and getting worse all the time
 - Transient faults:
 - Small (low charge) transistors are more easily flipped
 - Even low-energy particles can flip a bit now
 - Permanent faults:
 - Small transistors and wires deform and break more quickly
 - Higher temperatures accelerate the process
- Progression of transient faults
 - Memory (DRAM) was hit first: denser, smaller devices than SRAM
 - Then on-chip memory (SRAM)
 - Logic is starting to have problems...

Moore's Good Effect on Reliability

- The key to providing reliability is **redundancy**
 - The same scaling that makes devices less reliable...
 - Also increase device density to enable redundancy
- Examples
 - Error correcting code for memory (DRAM) and caches (SRAM)
 - Core-level redundancy: paired-execution, hot-spare, etc.
 - Intel's Core i7 (Nehalem) uses 8 transistor SRAM cells
 - Versus the standard 6 transistor cells
- Big open questions
 - Can we protect logic efficiently? (without 2x or 3x overhead)
 - Can architectural techniques help hardware reliability?
 - Can software techniques help?

Power Implications on Software

- Software-controlled dynamic voltage/frequency scaling
 - OS? Application?
 - Example: video decoding
 - Too high a clock frequency – wasted energy (battery life)
 - Too low a clock frequency – quality of video suffers
- Managing low-power modes
 - Don't want to “wake up” the processor every millisecond
- Tuning software
 - Faster algorithms can be converted to lower-power algorithms
 - Via dynamic voltage/frequency scaling
- Exploiting parallelism & heterogeneous cores
 - NVIDIA Tegra 3: 5 cores (4 “normal” cores & 1 “low power” core)
- Specialized hardware accelerators